

Human Pose Estimation for Video Game Control

Rakesh Johny, Tom Li, Aditya Narayanan, Albert Xia

Abstract

Current methods of interacting with computers are flawed in a couple of key ways: they fail to map physically-intuitive motions to their computer control counterparts, and they rely heavily on a user's fine-motor skills, which are heavily impacted by factors such as muscle coordination disabilities and old age. In this paper, we propose a system of computer control through gesture tracking via real-time human pose estimation on an embedded device, which is capable of addressing these issues, by mapping general physically-intuitive motions into computer control. Furthermore, our device boasts low latency in the human-computer interaction, and is minimally intrusive to the computer being controlled. We test the viability of our system as a computer-control device by playing two video games using only gestures.

1. Introduction

The human-computer interfaces behind many modern video games use either handheld joystick controllers or keyboard and mouse input, which are examples of typical human input device for computers, monitors, televisions, etc. There are two major drawbacks to traditional human input systems that are often overlooked: the input needed to achieve a certain effect has little to no correspondence to a physically-intuitive set of motions, and more importantly, inputs are heavily dependent on the user's fine motor skills. Fine motor skills used to interact with modern technology are heavily affected by factors such as muscular coordination disorders and old age. Computer vision, particularly the processes of human pose estimation and motion tracking may allow us to create human-computer interfaces that link computer control to physically-intuitive human motion inputs with little dependence on fine motor ability or additional physical input devices. We propose a vision-based computer-input system that maps gestures and motion of the user's body into computer input. We test the system's effectiveness with the playability of two simple video games as metrics. In order to demonstrate the versatility of our phased-approach to gesture-based control, we select a car-racing game and the Google Chrome Dino Racer game as

test video games. Such a system, when expanded to replicating general keyboard + mouse control, would be instrumental in improving the accessibility of modern technology, with implementations far beyond the scope of computer games.

2. Related Work

2.1. Human Pose Estimation

Some of the largest components of our system fall into the category of so called Human Pose Estimation or the ability to properly:

- Identify users as sources of input, from a video stream
- Segment the user's body to isolate relevant portions of data
- Compute the location, in image coordinates, of key-points on the user's body.

The area of Human Pose Estimation is widely studied. There are a wide array of robust, optimized solutions for gathering 2D image coordinates of keypoints from images of users.

2.1.1 Openpose

One popular solution to realtime human pose estimation, derived from [1], uses Part Affinity Fields (PAFs) to learn to associate body parts with individuals in the image, and open-source code for the so called openpose library shows promising results for real-time human body segmentation.




Figure 1. A demonstration of openpose performing human body segmentation.

2.1.2 MoveNet and PoseNet

PoseNet and its more recent counterpart MoveNet are fully convolutional neural network pose estimators built on existing ResNet or MobileNet backbones (citation here). Benchmarks of PoseNet and MoveNet show that it is viable to run either estimator in real time on embedded devices with appropriate model quantizations (citation here), with PoseNet being slightly faster when compared to MoveNet at identical levels of quantization. In comparison with Openpose, PoseNet/MoveNet are shown to perform slightly better on instance segmentation of the person class from the COCO dataset.

2.2. Gesture Identification/Motion Tracking

Openpose and MoveNet/PoseNet give us promising ways to identify locations of human body features (hands, arms, etc.) in a video stream. However, this is not enough to categorize a user's motion as a particular command to a computer system. To be able to extract feature motion over time as a useful input to our car racing video game, we will have to implement motion tracking over a video stream.

A number of guides exist for implementing feature tracking over video. [3] and [2] demonstrate simple object tracking across video frames, and [2] demonstrates a system shown to be accurate with motion of hands. We drew on [3] and [2] in developing a tracking algorithm for the identification of a human user ducking or jumping from a keypoint cloud.

3. Methodology

Our human computer interface consists of three major steps: Pose Estimation, or feature location, Feature Tracking and classification of motions, and video game interfacing. In addition, we adopt a multi-phased approach to our human computer interface. Phase 1 consists of a robust and general human pose estimator and keypoint detector. The work in Phase 1 is task/game agnostic, as we implement identical versions of this phase for the car-racer game as well as the chrome dinosaur game. Phase 2 is a game-specific keypoint classifier, where the output from Phase 1 is classified into a number of computer commands in a way dependent on the physical motions being replicated in the game as well as the computer commands the game is traditionally played with.

The motivation behind such a two phase approach rather than building a classifier that directly maps input video streams to computer commands is to improve the versatility of our human-computer interaction system. A fully general Phase 1 implemented independent of the game in question allows the system to be more easily adapted to any

particular set of computer commands. We demonstrate the versatility of this system through the implementation of our human-computer interface for two distinct video games that utilize identical Phase 1 code.

3.1. Testing Pose Estimation Methods

Before we develop phase 1 of our human-computer interface, we test a number of proposed solutions for human pose estimation to determine which ones are best-suited to our needs. The main criteria we require for our human pose estimator are:

- Real time operation. The need for this criterion is obvious, as we require low latency classification in order to mimic low-latency commands for video games. Essentially, the user should feel as if the delay between executing a physical command and the video game reflecting such a command is instantaneous.
- The system should be minimally intrusive on the computer used to play the game. This essentially means that we cannot use computing resources so significant that usability of the computer itself is diminished.

3.1.1 Openpose Python API

The openpose package offers a Python API for implementing inference over video feeds using one of the openpose models. We began testing the openpose Python API for our task, but found that installing the API resulted in dependency issues across multiple team members' computers and decided to instead test the openpose pretrained models implemented independently of the openpose python API.

3.1.2 Openpose model implemented in OpenCV

(citation) offers a method to perform inference on CPU/GPU using pretrained versions of the openpose hand/body pose estimators. For inference with the openpose pretrained model through OpenCV, we observed only 1-2 FPS on a 2.3 GHz Intel i9 CPU. We expect that running the same model with CUDA acceleration would result in true real-time performance, but as one of our design constraints for our system, we imposed that we cannot use significant GPU resources in order to maintain minimal intrusiveness onto the gaming machine.

3.1.3 PoseNet and the Jetson Nano

Due to the poor performance of the Openpose models, we began to look at different methods for real-time pose estimation. One option that was promising was to completely offload the task of human pose estimation to an edge-device to reduce the computational burden on the device used to

play the games. Such an embedded device would ideally be compact yet powerful enough to run quantized versions of pose estimators in real time. We found the NVIDIA Jetson Nano to be an ideal candidate for such an embedded device.

With a 128-core NVIDIA GPU on a single-board computer footprint, the Jetson Nano is a popular choice for deep model inference on embedded computers. Perhaps more important than the hardware itself, however, is the specific Tensor RT optimized models available for use on the Jetson Nano. Through a number of optimization techniques such as model quantization, unused output elimination and horizontal layer fusion/layer aggregation, the Tensorflow with Tensor RT (TF-TRT) backend is able to greatly improve inference time on the Jetson Nano. TF-TRT is able to output optimized models from Open Neural Network Exchange Format (ONNX) and for popular tasks such as human pose estimation, ONNX model formats are readily available from NVIDIA Developers.

3.2. Pose Estimation

Using a variant of PoseNet built on a ResNet-18 backbone, we observe 15-17 Frames Per Second (FPS) on the Jetson Nano. The so called Pose-ResNet-18-body outputs 17 detected keypoints in image coordinate format. 15-17 FPS performance is smooth enough keypoint detection to allow us to play these simple games in real time.

3.3. Feature Tracking and Classification

3.3.1 Feature Classification For Steering

3.3.2 Feature Classification For Jumping/Ducking

3.4. Game Interface

4. Results

5. Conclusion and Future Works

References

- [1] Z. Cao, G. Hidalgo Martinez, T. Simon, S. Wei, and Y. A. Sheikh. Openpose: Realtime multi-person 2d pose estimation using part affinity fields. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 2019.
- [2] Charlie Liu. Gesture detection with a raspberry pi. 2017.
- [3] Adrian Rosebrock. Simple object tracking with opencv. 2018.